

Sentiment Prediction based on Drug reviews

Team: Gargeyi Baipa , Naveen Donthineni

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
import os
!pip install wget
!pip install wordcloud
from wordcloud import WordCloud
!pip install yellowbrick
```

Requirement already satisfied: wget in c:\programdata\anaconda3\lib\site-packages (3.2)

Requirement already satisfied: wordcloud in c:\programdata\anaconda3\lib\site-packages (1.8.1)

Requirement already satisfied: matplotlib in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (3.1.0)

Requirement already satisfied: numpy>=1.6.1 in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (1.16.4)

Requirement already satisfied: pillow in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (6.0.0)

Requirement already satisfied: cycler>=0.10 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (0.10.0)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (1.1.0)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.4.0)

Requirement already satisfied: python-dateutil>=2.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.8.0)

Requirement already satisfied: six in c:\programdata\anaconda3\lib\site-packages (from cycler>=0.10->matplotlib->wordcloud) (1.12.0)

Requirement already satisfied: setuptools in c:\programdata\anaconda3\lib\site-packages (from kiwisolver>=1.0.1->matplotlib->wordcloud) (41.0.1)

Requirement already satisfied: yellowbrick in c:\programdata\anaconda3\lib\site-packages (1.2)

Requirement already satisfied: scikit-learn>=0.20 in c:\users\adari\appdata\roaming\python\python37\site-packages (from yellowbrick) (0.23.2)

Requirement already satisfied: matplotlib!=3.0.0,>=2.0.2 in c:\programdata\anaconda3\lib\site-packages (from yellowbrick) (3.1.0)

Requirement already satisfied: numpy>=1.13.0 in c:\programdata\anaconda3\lib\site-packages (from yellowbrick) (1.16.4)

Requirement already satisfied: scipy>=1.0.0 in c:\programdata\anaconda3\lib\site-packages (from yellowbrick) (1.2.1)

Requirement already satisfied: cycler>=0.10.0 in c:\programdata\anaconda3\lib\site-packages (from yellowbrick) (0.10.0)

Requirement already satisfied: joblib>=0.11 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn>=0.20->yellowbrick) (0.13.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn>=0.20->yellowbrick) (2.1.0)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (1.1.0)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (2.4.0)

Requirement already satisfied: python-dateutil>=2.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib!=3.0.0,>=2.0.2->yellowbrick) (2.8.0)

Requirement already satisfied: six in c:\programdata\anaconda3\lib\site-packages (from cycler>=0.10.0->yellowbrick) (1.12.0)

Requirement already satisfied: setuptools in c:\programdata\anaconda3\lib\site-packages (from kiwisolver>=1.0.1->matplotlib!=3.0.0,>=2.0.2->yellowbrick) (41.0.1)

```
In [91]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import seaborn as sn
import matplotlib.pyplot as plt
from yellowbrick.classifier import ROCAUC
from yellowbrick.datasets import load_spam
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from yellowbrick.classifier import ClassificationReport
```

Import the data file for the analysis

```
In [189]: import wget
```

```
In [190]: import re
```

```
In [191]: wget.download('https://archive.ics.uci.edu/ml/machine-learning-databases/0046
2/drugsCom_raw.zip')
```

```
Out[191]: 'drugsCom_raw (23).zip'
```

```
In [193]: import zipfile
```

```
In [192]: with zipfile.ZipFile("drugsCom_raw.zip","r") as zip_ref:
zip_ref.extractall()
```

```
In [194]: dftrain = pd.read_csv('drugsComTest_raw.tsv',delimiter = '\t')
```

```
In [195]: dftrain = pd.read_csv('drugsComTrain_raw.tsv',delimiter = '\t')
```

```
In [196]: combine = [dftrain,dftrain]
df = pd.concat(combine)
```

In [197]: `df.head()`

Out[197]:

	Unnamed: 0	drugName	condition	review	rating	date	usefulCount
0	206461	Valsartan	Left Ventricular Dysfunction	"It has no side effect, I take it in combinati...	9.0	May 20, 2012	27
1	95260	Guanfacine	ADHD	"My son is halfway through his fourth week of ...	8.0	April 27, 2010	192
2	92703	Lybrel	Birth Control	"I used to take another oral contraceptive, wh...	5.0	December 14, 2009	17
3	138000	Ortho Evra	Birth Control	"This is my first time using any form of birth...	8.0	November 3, 2015	10
4	35696	Buprenorphine / naloxone	Opiate Dependence	"Suboxone has completely turned my life around...	9.0	November 27, 2016	37

Data Cleaning

In [198]: `df['date'] = pd.to_datetime(df['date'])`
`dftrain['date'] = pd.to_datetime(dftrain['date'])`

In [159]: `df_drug_review=pd.DataFrame(dftrain)`
`df_drug_review.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 161297 entries, 0 to 161296
Data columns (total 7 columns):
Unnamed: 0      161297 non-null int64
drugName        161297 non-null object
condition        160398 non-null object
review           161297 non-null object
rating           161297 non-null float64
date             161297 non-null object
usefulCount      161297 non-null int64
dtypes: float64(1), int64(2), object(4)
memory usage: 8.6+ MB
```

There are 161297 records exists in the dataset and 6 variablea are available

Look at the summary stastics of numerical fileds

```
In [160]: df_drug_review.describe()[['rating', 'usefulCount']]
```

```
Out[160]:
```

	rating	usefulCount
count	161297.000000	161297.000000
mean	6.994377	28.004755
std	3.272329	36.403742
min	1.000000	0.000000
25%	5.000000	6.000000
50%	8.000000	16.000000
75%	10.000000	36.000000
max	10.000000	1291.000000

As we see, mean rating value of all drugs is 6.9 and around average of 28 people found the drug reviews are useful

Check the dataset for the missing values

```
In [162]: df_drug_review.isnull().values.any()
```

```
Out[162]: True
```

```
In [163]: df_drug_review.columns[df_drug_review.isnull().any()].tolist()
```

```
Out[163]: ['condition']
```

Looks like "condition" attribute is missing the values but rest of the fields are ok

```
In [164]: size = df_drug_review.shape[0]

print ("Total Size of the dataset : ", size)

total_na = df_drug_review.isnull().sum(axis = 0)['condition']
print ("Null values : ", total_na)

print ("PERCENTAGE : ", (total_na/size)*100)
```

```
Total Size of the dataset : 161297
Null values : 899
PERCENTAGE : 0.5573569254232875
```

"condition" variable is having missing the values and % of missing values are 0.5%. So we have removed them from the analysis

Check for Duplicate values

```
In [165]: df_drug_review.duplicated()
```



```
Out[165]: 0      False
          1      False
          2      False
          3      False
          4      False
          5      False
          6      False
          7      False
          8      False
          9      False
         10      False
         11      False
         12      False
         13      False
         14      False
         15      False
         16      False
         17      False
         18      False
         19      False
         20      False
         21      False
         22      False
         23      False
         24      False
         25      False
         26      False
         27      False
         28      False
         29      False
          ...
        161267  False
        161268  False
        161269  False
        161270  False
        161271  False
        161272  False
        161273  False
        161274  False
        161275  False
        161276  False
        161277  False
        161278  False
        161279  False
        161280  False
        161281  False
        161282  False
        161283  False
        161284  False
        161285  False
        161286  False
        161287  False
        161288  False
        161289  False
        161290  False
        161291  False
        161292  False
```

```
161293    False
161294    False
161295    False
161296    False
Length: 161297, dtype: bool
```

Convert the all column names into lower case

```
In [166]: df_drug_review.columns=df_drug_review.columns.str.lower()
```

How many unique drugs and conditions exists in the dataset?

```
In [167]: ##number of unique drugs present in the dataset
df_drug_review['drugname'].nunique()
```

```
Out[167]: 3436
```

```
In [168]: ##number of unique conditions present in the dataset
df_drug_review['condition'].nunique()
```

```
Out[168]: 884
```

There are 3431 drugs and 884 unique health conditions exists in the dataset

Exploratory Data Analysis (EDA)

Let us analyse the drugs attribute

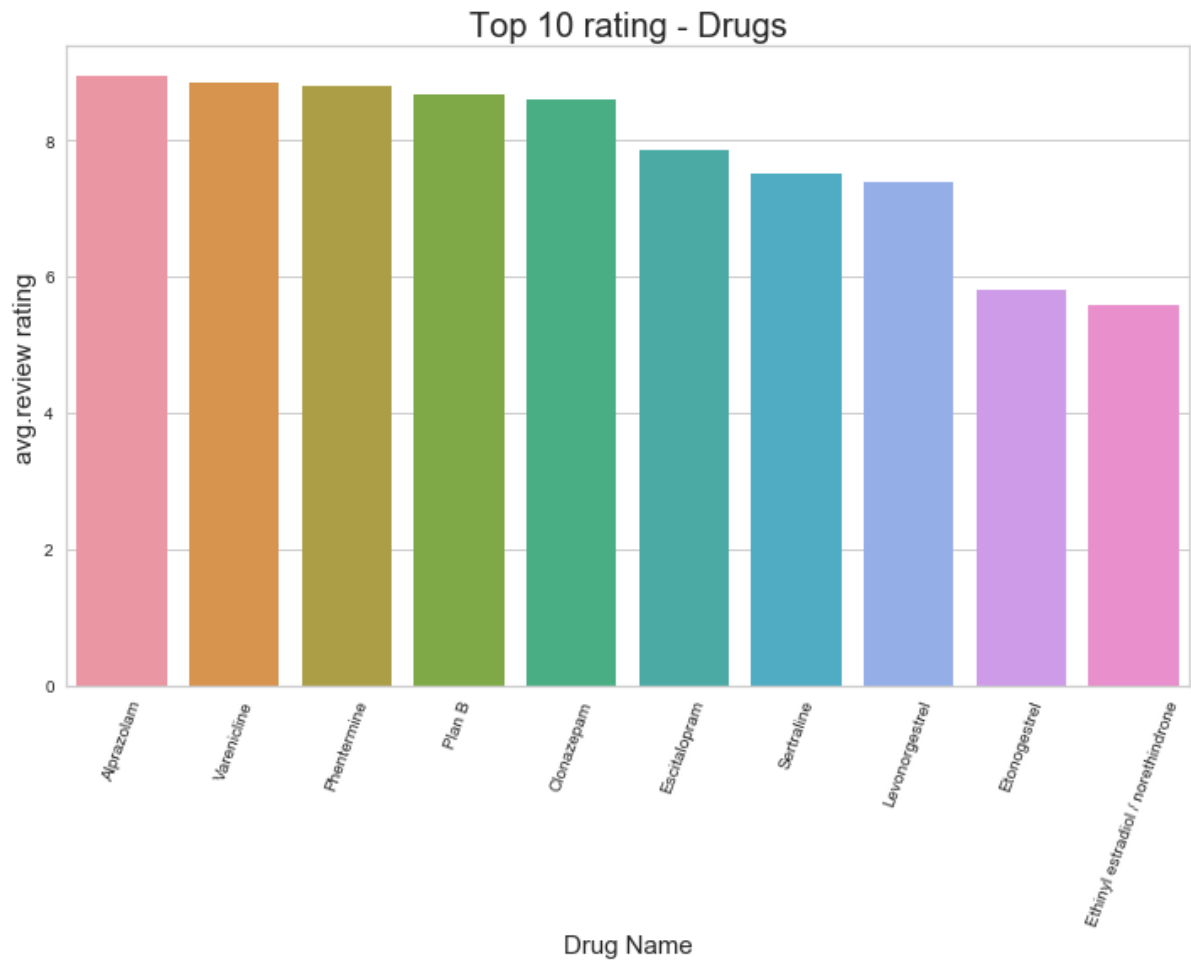
What are Top 10 drugs in the dataset?

```
In [169]: ##Find list of drugs which are having rating 10
df_drug_review_10 = dict(df_drug_review.loc[df_drug_review.rating == 10, "drug
name"].value_counts())
drugname = list(df_drug_review_10.keys())
drug_rating = list(df_drug_review_10.values())
df_drug_review_10=df_drug_review.groupby(['drugname']).count()['rating'].reset
_index().sort_values(by='rating', ascending=False)
#Calculate the mean review rating for Top 10 drugs
top_10_drugs=pd.DataFrame(drugname[0:10],columns=['drugname'])
top_10_drugs=pd.merge(top_10_drugs,df_drug_review,how='inner', left_on=['drugn
ame'], right_on = ['drugname'])
top_10_drugs_avg=top_10_drugs.groupby(['drugname']).mean()['rating'].reset_ind
ex()
top_10_drugs_avg=top_10_drugs_avg.sort_values(by='rating', ascending=False)

##plot top 10 drugs and avg.review rating

drugname_avg = list(top_10_drugs_avg.drugname)
drug_rating_avg = list(top_10_drugs_avg.rating)
plt.figure(figsize=(10,6))
plt.xticks(rotation=70)
plt.tight_layout()
plt_rating = sns.barplot(x = drugname_avg, y =drug_rating_avg)
plt.xlabel("Drug Name", size=15)
plt.ylabel("avg.review rating", size=15)
plt.title("Top 10 rating - Drugs",size=20)
```

Out[169]: Text(0.5, 1, 'Top 10 rating - Drugs')



```
In [170]: #merge datasets to get combined view of avg.ratings and review counts
top_10_drugs_avg=pd.merge(top_10_drugs_avg,df_drug_review_10,how='inner', left
_on=['drugname'], right_on = ['drugname'])
#top_10_drugs_avg.head(10)
```

Alprazolam,Varencline,phentermine,Plan B,Clonazephram,Escitalopram,Sertaline are some of top 10 drugs as per the avg. review ratings

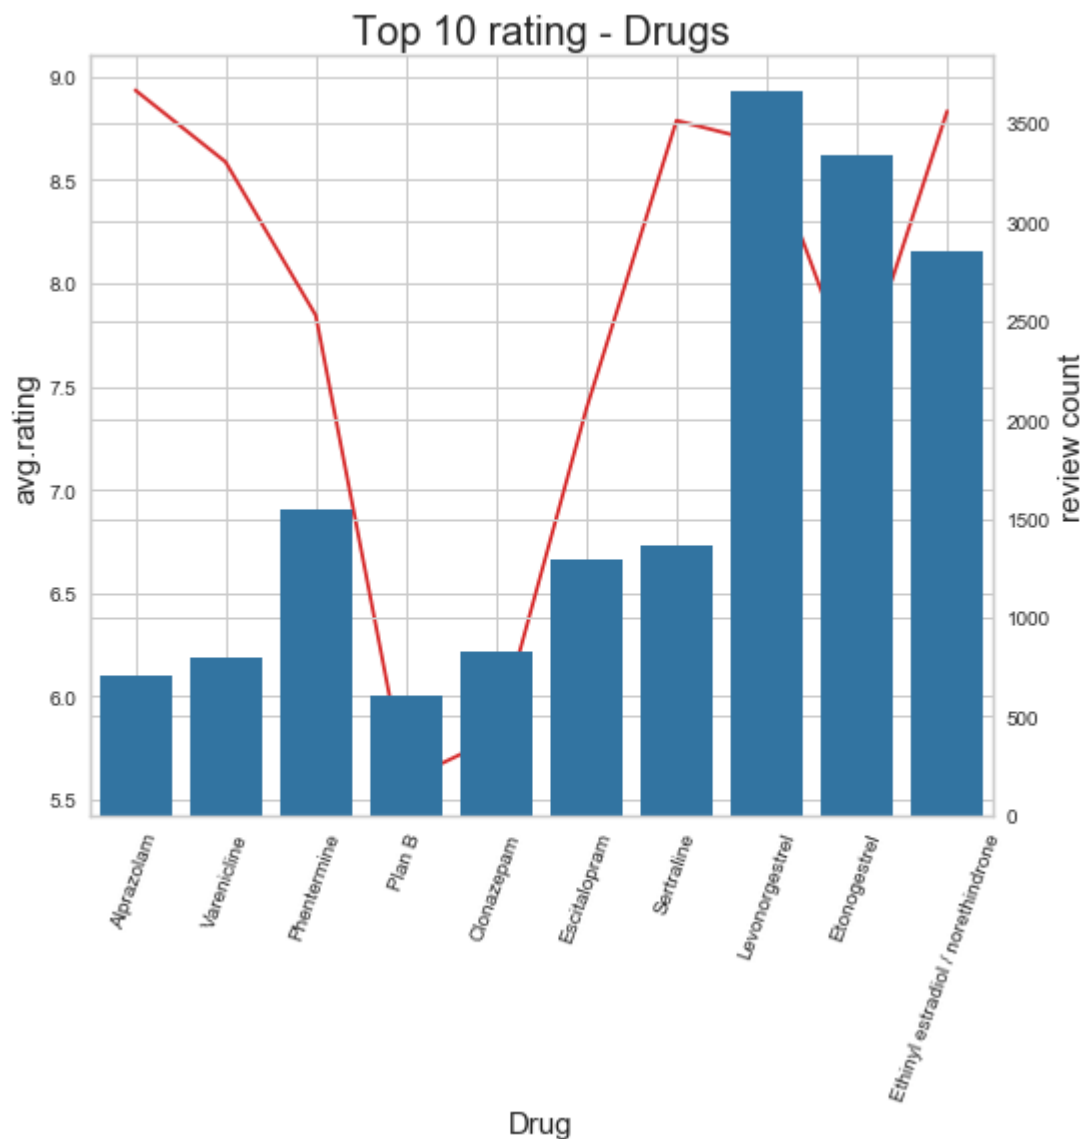
What are the review counts for the top 10 drugs?

```
In [171]: #plot the line chart for review count

plt.figure(figsize=(7,6))
plt.xticks(rotation=70)
color1 = 'tab:red'
plt.tight_layout()
plt_rating = sns.lineplot(x='drugname', y='rating_x', data = top_10_drugs_avg,
color=color1)
plt.xlabel("Drug", size=15)
plt.ylabel("avg.rating", size=15)
plt.title("Top 10 rating - Drugs",size=20)

#line chart for avg ratings for top 10 drugs

ax2 = plt_rating.twinx()
color = 'tab:blue'
#line plot creation
ax2.set_ylabel('Avg rating', fontsize=16)
ax2 = sns.barplot(x='drugname', y='rating_y', data = top_10_drugs_avg, color=
color)
plt.ylabel("review count", size=15)
#show plot
plt.show()
```



Etonogestrel drug is having highest number of review counts, review rating is low on the other hand Alprazolam drug is having highest review rating, less review count, this is an intuitive insight

Which Drug reviews are found useful?

```
In [172]: ##Look at the top 10 drugs which are having highest useful counts
top_10_drugs_usefulcount=top_10_drugs.groupby(['drugname']).sum()['usefulcount'].reset_index().sort_values(by='usefulcount', ascending=False)
top_10_drugs_count_avgrating=top_10_drugs.groupby(['drugname']).mean()['rating'].reset_index().sort_values(by='rating', ascending=False)
drugs_count_avgrating=pd.merge(top_10_drugs_usefulcount,top_10_drugs_count_avgrating,how='inner', left_on=['drugname'], right_on = ['drugname'])
drugs_count_avgrating=drugs_count_avgrating.sort_values(by='usefulcount', ascending=False)
```

```
In [173]: drugs_count_avgrating.head(10)
```

```
Out[173]:
```

	drugname	usefulcount	rating
0	Sertraline	58154	7.497794
1	Escitalopram	52280	7.843653
2	Phentermine	46494	8.786131
3	Levonorgestrel	36769	7.391031
4	Clonazepam	32339	8.586041
5	Alprazolam	29826	8.933144
6	Etonogestrel	21917	5.802458
7	Ethinyl estradiol / norethindrone	21510	5.591930
8	Varenicline	15184	8.831234
9	Plan B	7322	8.674917


```

In [174]: drugname_count = list(drugs_count_avgrating.drugname)
drug_usefulcount = list(drugs_count_avgrating.usefulcount)
plt.figure(figsize=(7,6))
plt.xticks(rotation=70)
plt.tight_layout()
color1 = 'tab:blue'
plt_rating = sns.barplot(x='drugname', y='usefulcount', data = drugs_count_avgrating, color=color1)
plt.xlabel("Drug", size=15)
plt.ylabel("usefulcount", size=15)
plt.title("Top 10 drugs useful reviews and avg.ratings",size=20)

#line chart for avg ratings for top 10 drugs

ax2 = plt_rating.twinx()
color = 'tab:red'
#line plot creation
ax2.set_ylabel('Avg rating', fontsize=16)
ax2 = sns.lineplot(x='drugname', y='rating', data = drugs_count_avgrating, color=color)
plt.ylabel("avg rating", size=15)
#show plot
plt.show()

```




Intersting Insight : People are found some of the drug reviews are useful eventhough thier avg.review ratings are less.As we see,"Sertaline" drug reviews are found useful eventhough it has avg.review rating between 7-7.5 .But this scnerio is different for the drugs which are having highest avg.review ratings ex:Alprazolam

What are the Top 10 conditions in the dataset?

```
In [175]: ## Find out top 10 conditions
df_condition=df_drug_review.groupby(['condition']).count()['rating'].reset_index().sort_values(by='rating',ascending=False)
df_condition=df_condition[0:10]
df_condition

## Plot top 10 conditions
plt.figure(figsize=(10,6))
plt.xticks(rotation=70)
plt.tight_layout()
chart=sns.barplot(x="rating", y="condition",
                  data=df_condition,order=df_condition.sort_values('rating',ascending=False).condition,color='grey')
plt.xlabel("count of drugs", size=15)
plt.ylabel("conditions", size=15)
plt.title("Top 10 Conditions",size=20)
```

Out[175]: Text(0.5, 1, 'Top 10 Conditions')



Birth Control,Depression and Pain are the top 3 health conditions out of top 10.

What is the avg drug review rating and number of drugs available for each top 10 health conditions?

```
In [176]: df_drug_condition=df_drug_review.groupby(['condition'])['drugname'].nunique().  
reset_index()
```

```

In [177]: ## Find out top 10 conditions and avg drug ratings for these conditions
df_condition_avgrating=df_drug_review.groupby(['condition']).mean()['rating'].
reset_index().sort_values(by='rating',ascending=False)
df_condition_avgrating=pd.merge(df_condition_avgrating,df_condition,how='inner',
left_on=['condition'], right_on = ['condition'])

##find the drug counts for each condition

condition_drug_count=pd.merge(df_condition,df_drug_condition,how='inner', left
_on=['condition'], right_on = ['condition'])
condition_drug_count=condition_drug_count.rename(columns={"rating": "count"})
#print(condition_drug_count)


##combined avg rating and drug count

df_condition_avgrating=pd.merge(df_condition_avgrating,condition_drug_count,ho
w='inner', left_on=['condition'], right_on = ['condition'])
df_condition_avgrating.head(10)
df_condition_avgrating.rename(columns={"rating_x":"avg_rating","rating_y":"cou
nt_y","drugname":"drug_count"},inplace=True)
## Plot top 10 conditions anf avg.review ratings
topcondition_count = list(df_condition_avgrating.condition)
condition_avgrating = list(df_condition_avgrating.avg_rating)
plt.figure(figsize=(10,6))
plt.xticks(rotation=70)
plt.tight_layout()
color = 'tab:blue'
plt_rating = sns.barplot(x='condition', y='avg_rating', data = df_condition_av
grating,color=color)
plt.xlabel("health condition", size=10)
plt.ylabel("avg.rating", size=10)
#plt.title("Top 10 conditions avgreviews and drug counts",size=20)

#line chart for avg ratings for top 10 drugs

ax2 = plt_rating.twinx()
color = 'tab:red'
#line plot creation
ax2.set_ylabel('drug count', fontsize=14)
ax2 = sns.lineplot(x='condition', y='drug_count', data = df_condition_avgratin
g,color=color)
plt.ylabel("drug count", size=10)
#show plot
plt.show()

```



Intersting Insight 2: Some of the top health conditions avg review rating is high eventhough very few drugs available,ex:Weight loss,Obesity.At the same time,some health conditions are having less avg review rating but available drugs are very high example: Birth control

```


In [178]: ## Find drugs ratings for top 10 health conditions
condition_drug3_count=pd.merge(df_condition,df_drug_review,how='inner', left_on=['condition'], right_on = ['condition'])
condition_drug3_count=condition_drug3_count.rename(columns={"rating_y": "rating"})
df_condition_top3_avgrating=condition_drug3_count.groupby(['condition','drugname']).mean()['rating'].reset_index().sort_values(by='rating',ascending=False)

##Plot the ratings for all top 10 conditions and see how many drugs are having highest ratings and lowest ratings for each condition

plt.figure(figsize=(15,8))
plt.xticks(rotation=70)
plt_box=sns.boxplot(y='rating', x='condition',
                    data=df_condition_top3_avgrating,
                    width=0.5,
                    palette="colorblind")
plt_box=sns.stripplot(y='rating', x='condition',
                     data=df_condition_top3_avgrating,
                     jitter=True,
                     marker='o',
                     alpha=0.5,
                     color='black')
plt.xlabel("condition", size=15)
plt.ylabel("rating", size=15)
#plt.title("Top 10 Conditions - Drugs and ratings",size=20)

```

Out[178]: Text(0, 0.5, 'rating')



Insight 3: Birth control conditions is having around 170 drugs but small number of drugs are having highest reviews but in case of weigh all the drugs are having very good review ratings

What are the bottom 10 health conditions?

```
In [179]: ## Find out bottom 10 conditions
df_condition_bottom=df_drug_review.groupby(['condition']).count()['rating'].re
set_index().sort_values(by='rating',ascending=False)
df_condition_bottom=df_condition_bottom.tail(10)
df_condition_bottom
##find the drug counts for each condition

bottomcondition_drug_count=pd.merge(df_condition_bottom,df_drug_condition,how=
'inner', left_on=['condition'], right_on = ['condition'])
bottomcondition_drug_count=bottomcondition_drug_count.rename(columns={"rating"
: "count"})
bottomcondition_drug_count

##combined avg rating and drug count

df_condition_bottom_avgrating=pd.merge(df_condition_bottom,bottomcondition_dru
g_count,how='inner', left_on=['condition'], right_on = ['condition'])
df_condition_bottom_avgrating.head(10)
df_condition_bottom_avgrating.rename(columns={"rating_x":"avg_rating","rating_
y":"count_y","drugname":"drug_count"},inplace=True)
df_condition_bottom_avgrating
```

Out[179]:

	condition	rating	count	drug_count
0	Syringomyelia	1	1	1
1	Systemic Candidiasis	1	1	1
2	95 users found this comment helpful.	1	1	1
3	Cerebrovascular Insufficiency	1	1	1
4	Ramsay Hunt Syndrome	1	1	1
5	Radionuclide Myocardial Perfusion Study	1	1	1
6	Rabies Prophylaxis	1	1	1
7	Q Feve	1	1	1
8	Bartonellosis	1	1	1
9	Anti NMDA Receptor Encephalitis	1	1	1

Syringomyelia, Systemic Candidiasis, Cerebrovascular Insufficiency, Ramsay Hunt Syndrome, Radionuclide Myocardial Perfusion Study, Rabies Prophylaxis, Q Feve, Bartonellosis, Anti NMDA Receptor Encephalitis are bottom 10 health conditons.

Insight 4: All bottom health conditions are having only one drug and review count available.

Time analysis

```
In [180]: # create new data frame for time analysis
drug_review_date=df_drug_review[["drugname","date", "rating","usefulcount"]]
#add month variable to the dataframe
drug_review_date['month'] = pd.DatetimeIndex(drug_review_date['date']).month
#add year variable to the dataframe
drug_review_date['year'] = pd.DatetimeIndex(drug_review_date['date']).year
#convert year variable to string for charing purpose
drug_review_date['year'] = drug_review_date['year'].apply(str)

#create data frame to see the data by year

drug_review_year=drug_review_date.groupby(['year']).count()['rating'].reset_index()

drug_review_year['year'] = drug_review_year['year'].apply(str)
drug_review_year=drug_review_year.set_index('year')
drug_review_year
```

```
C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:4: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

after removing the cwd from sys.path.

```
C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:6: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:8: SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```


See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

Out[180]:

	rating
year	
2008	5137
2009	11636
2010	8413
2011	11682
2012	10007
2013	12278
2014	12051
2015	27164
2016	34842
2017	28087

How the review rating counts changed year over year?


```
In [181]: #plot the review count over the years
ax = drug_review_year.plot(figsize=(12,6), color='r', marker='o')
ax.set_yticklabels(pd.Series(ax.get_yticks()).map(lambda x: '{:.0f}'.format(x)))
#ax.set_title("Ratings count by year",size=20)
ax.set_xlabel("Year",size=10)
ax.set_ylabel("review count",size=10)
ax.grid(color='grey', linestyle='-', linewidth=0.2)
```




Insight 5: As we see, reviews for the drugs gradually increased from 2010 to 2014 with few fluctuations and peaked in the year of 2016 before it sees sudden drop of review count in 2017.

How useful counts were changed year over year?

```
In [182]: ##create new dataframe to get the YOY useful counts
drug_usefulcount_year=drug_review_date.groupby(['year']).sum()['usefulcount'].
reset_index()
drug_usefulcount_year['year'] = drug_usefulcount_year['year'].apply(str)
drug_usefulcount_year=drug_usefulcount_year.set_index('year')

#plot the useful count over the years
ax = drug_usefulcount_year.plot(figsize=(12,6), color='r', marker='o')
ax.set_yticklabels(pd.Series(ax.get_yticks()).map(lambda x: '{:.0f}'.format(x
)))
#ax.set_title("usefulcount by year",size=20)
ax.set_xlabel("Year",size=10)
ax.set_ylabel("count",size=10)
ax.grid(color='grey', linestyle='-', linewidth=0.2)
```



Insight 6 : Number of people gradually increased who found drug reviews are useful from 2010 to 2014 and this count is peaked in 2015 before it sees sudden drop in subsequent years..Looks like, number of reviews for the drugs are decreased from 2016 so usage of drug reviews also decreased

How top 10 drugs ratings changed year over year?


```
In [183]: ##create new dataframe to answer this question
drug_review_year=drug_review_date.groupby(["drugname", 'year']).mean()['rating']
].reset_index()
top_10_drugs_avg=top_10_drugs_avg['drugname']
drug_review_year=pd.merge(top_10_drugs_avg,drug_review_year,how='inner', left_
on=['drugname'], right_on = ['drugname'])

##display sample records of this dataframe
drug_review_year.head(120)

##Plot the ratings for all top 10 conditions and see how many drugs are having
highest ratings and lowest ratings for each condition

plt.figure(figsize=(15,8))
plt.xticks(rotation=70)
plt_box=sns.lineplot(data=drug_review_year, x="year", y="rating", hue="drugnam
e")
plt.xlabel("year", size=15)
plt.ylabel("avg.rating", size=15)
plt.title("Top 10 drugs Change of reviews over the years",size=20)
plt_box.legend(loc='upper right', bbox_to_anchor=(1.4, 1))
```

Out[183]: <matplotlib.legend.Legend at 0x20451b20cf8>




Insight 7: ratings for the top 10 drugs from 2014-2015 and gradually decreased in the following years.Ex:ethinyl estradiol,etc

What is the correlation between ratings and useful count?

```
In [184]: corr = df_drug_review.corr()
# Plot heatmap to show correlation

plt.figure(figsize=(15,8))
plt.xticks(rotation=70)
plt_box=sns.heatmap(corr, annot=True)
```



Insight 8: correlation between rating and useful count is positive but it is very small i.e 0.23..we same same insight in top drugs avg.review and useful counts.

In [185]: `!pip install wordcloud`

```
Requirement already satisfied: wordcloud in c:\programdata\anaconda3\lib\site-packages (1.8.1)
Requirement already satisfied: pillow in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (6.0.0)
Requirement already satisfied: numpy>=1.6.1 in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (1.16.4)
Requirement already satisfied: matplotlib in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (3.1.0)
Requirement already satisfied: cyclor>=0.10 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (1.1.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.4.0)
Requirement already satisfied: python-dateutil>=2.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.8.0)
Requirement already satisfied: six in c:\programdata\anaconda3\lib\site-packages (from cyclor>=0.10->matplotlib->wordcloud) (1.12.0)
Requirement already satisfied: setuptools in c:\programdata\anaconda3\lib\site-packages (from kiwisolver>=1.0.1->matplotlib->wordcloud) (41.0.1)
```

In [186]: `from wordcloud import WordCloud`

What are the most influential words for top 10 drugs?

[illegible]

What are the most influential words for top 10 health conditions?

[illegible]

Sentiment Analysis

```
Requirement already satisfied: nltk in c:\programdata\anaconda3\lib\site-pack
ages (3.4.1)
Requirement already satisfied: six in c:\programdata\anaconda3\lib\site-packa
ges (from nltk) (1.12.0)
```

```
In [13]: import nltk
nltk.download(['punkt', 'stopwords'])

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\adari\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\adari\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!

Out[13]: True
```

The attribute "review" is analyzed below in greater detail

Count the words in each review

```
In [14]: df['wordcount'] = df['review'].apply(lambda x: len(x.split()))
```

Count the characters in each review

```
In [15]: df['charcount'] = df['review'].apply(lambda x: len(x))
```

Convert all the words in the review to lower case

```
In [16]: df['lowercase'] = df['review'].apply(lambda x: " ".join(word.lower() for word
in x.split()))
```

Remove all the punctuation

```
In [17]: df['punctuation'] = df['lowercase'].str.replace('[^\w\s]', '')
```

Porter stemmer is installed however the approach to removed stop words and lemmatize is expected to be a better approach for text analysis


```
In [18]: from nltk.stem import PorterStemmer
         from nltk.corpus import stopwords
         nltk.download('wordnet')
```

```
[nltk_data] Downloading package wordnet to
[nltk_data]   C:\Users\adari\AppData\Roaming\nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

```
Out[18]: True
```

Remove the stop words in english

```
In [19]: stop_words = stopwords.words('english')
```

```
In [20]: df['stopwords'] = df['punctuation'].apply(lambda x: " ".join(word for word in
         x.split() if word not in stop_words))
```

```
In [21]: #pd.Series(" ".join(df['stopwords']).split()).value_counts()[45:100]
pd.Series(" ".join(df['stopwords']).split()).value_counts()[0:45]
```

```
Out[21]: day                74226
i039m                71694
side                 69978
taking               68233
years                61929
i039ve              61260
first                59838
pain                 59660
effects             58930
take                 58291
months              56681
started             56428
like                 55493
get                  55013
days                53468
time                 51291
feel                 47361
2                    44971
would                44305
back                 41324
pill                 40974
3                    40334
one                  39930
weeks                39457
it039s              38813
week                 37749
took                 37379
also                 37242
period              36779
medication           36742
doctor              36573
weight              36401
got                  36145
month                34800
since                33729
life                 33212
don039t             32544
still                32459
bad                  32430
really              31692
much                 31142
anxiety              30593
never                29855
went                 29120
better               29049
dtype: int64
```

Create additional wrds that can be cleaned

```
In [22]: other_stopwords = ['don039t', 'it039s', '2', '3', 'get', 'would', 'i039m', 'i039ve',
'4', '5', '10', '6', 'two', 'getting', 'didn039t', 'also']
```

```
In [23]: df['clean_review'] = df['stopwords'].apply(lambda x: " ".join(word for word in
x.split() if word not in other_stopwords))
```

```
In [24]: df.dropna().shape
```

```
Out[24]: (213869, 13)
```

Lemmatization is performed using text blob

```
In [25]: !pip install -U textblob
```

```
Requirement already up-to-date: textblob in c:\programdata\anaconda3\lib\site
-packages (0.15.3)
Requirement already satisfied, skipping upgrade: nltk>=3.1 in c:\programdata
\anaconda3\lib\site-packages (from textblob) (3.4.1)
Requirement already satisfied, skipping upgrade: six in c:\programdata\anacon
da3\lib\site-packages (from nltk>=3.1->textblob) (1.12.0)
```

```
In [26]: #import text blob
from textblob import Word
```

```
In [27]: df['lemmatize'] = df['clean_review'].apply(lambda x: " ".join(Word(word).lemmat
ize() for word in x.split()))
```

calculate the sentiment ('Polarity') and subjectivity of each review)

```
In [28]: from textblob import TextBlob
```

```
In [29]: df['sentiment'] = df['lemmatize'].apply(lambda x: TextBlob(x).sentiment[0])
```

```
In [30]: df['subjectivity'] = df['lemmatize'].apply(lambda x: TextBlob(x).sentiment[1])
```

In [31]: `df.head()`

Out[31]:

	Unnamed: 0	drugName	condition	review	rating	date	usefulCount	wordcount
0	206461	Valsartan	Left Ventricular Dysfunction	"It has no side effect, I take it in combinati...	9.0	May 20, 2012	27	1
1	95260	Guanfacine	ADHD	"My son is halfway through his fourth week of ...	8.0	April 27, 2010	192	14
2	92703	Lybrel	Birth Control	"I used to take another oral contraceptive, wh...	5.0	December 14, 2009	17	13
3	138000	Ortho Evra	Birth Control	"This is my first time using any form of birth...	8.0	November 3, 2015	10	8
4	35696	Buprenorphine / naloxone	Opiate Dependence	"Suboxone has completely turned my life around...	9.0	November 27, 2016	37	12

In [32]: `df.drop(['date', 'lowercase', 'punctuation', 'stopwords', 'clean_review'], axis = 1, inplace = True)`

In [33]: `df.head()`

Out[33]:

	Unnamed: 0	drugName	condition	review	rating	usefulCount	wordcount	charcou
0	206461	Valsartan	Left Ventricular Dysfunction	"It has no side effect, I take it in combinati...	9.0	27	17	;
1	95260	Guanfacine	ADHD	"My son is halfway through his fourth week of ...	8.0	192	141	7.
2	92703	Lybrel	Birth Control	"I used to take another oral contraceptive, wh...	5.0	17	134	7!
3	138000	Ortho Evra	Birth Control	"This is my first time using any form of birth...	8.0	10	89	4.
4	35696	Buprenorphine / naloxone	Opiate Dependence	"Suboxone has completely turned my life around...	9.0	37	124	7.

In []: `dfsent = df.sort_values(by = 'sentiment')`

Create a class label based on sentiment

In [34]: `df["class_label"] = df["sentiment"].apply(lambda x: 'neutral' if x==0 else ('negative' if x<0 else 'positive'))`

In [35]: `df.tail()`


Out[35]:

	Unnamed: 0	drugName	condition	review	rating	usefulCount	wordcount	ch
53761	159999	Tamoxifen	Breast Cancer, Prevention	"I have taken Tamoxifen for 5 years. Side effe...	10.0	43	97	
53762	140714	Escitalopram	Anxiety	"I've been taking Lexapro (escitalopram)...	9.0	11	130	
53763	130945	Levonorgestrel	Birth Control	"I'm married, 34 years old and I have no ...	8.0	7	149	
53764	47656	Tapentadol	Pain	"I was prescribed Nucynta for severe neck/shou...	1.0	20	34	
53765	113712	Arthrotec	Sciatica	"It works!!!"	9.0	46	2	

Correlation plot between sentiment and the rating

In [36]: `import seaborn as sns`
`plt = sns.lineplot(data=df, x='rating', y='sentiment')`
`plt`

Out[36]: `<matplotlib.axes._subplots.AxesSubplot at 0x20454231da0>`



We can see that there is positive correlation, as the rating increases the sentiment is also increased. Showing the positive sentiment associated with higher rating

```
In [83]: df2 = df.iloc[:,10:12]
```

```
In [84]: df2.head()
```

```
Out[84]:
```

	subjectivity	class_label
0	0.000000	neutral
1	0.482812	positive
2	0.367599	positive
3	0.618750	positive
4	0.384359	positive

```
In [85]: df2.groupby('class_label').size()
```

```
Out[85]: class_label
negative      70105
neutral       11661
positive     133297
dtype: int64
```

```
In [ ]: df2.groupby('class_label').count().plot.bar()
```

negative reviews account to 32 %,positive 61.9% and neutral reviews to 5 % of the data

Calculate term frequency and inverse document frequency of the review that is lemmatized using `TfidfVectorizer`

```
In [87]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [90]: v = TfidfVectorizer()
attribute = v.fit_transform(df['lemmatize'])
claslabel = df.class_label
#print(claslabel)
```

```
In [ ]: attribute
```

```
In [ ]: df.head
```

Machine learning to predict sentiment based on review

1.Logistic Regression

```
In [92]: x_train,x_test,y_train,y_test = train_test_split(attribute,claslabel,test_size  
=0.30,random_state=0)
```



```
In [97]: ##Build logistic regression
from yellowbrick.classifier import ConfusionMatrix
model = LogisticRegression(multi_class='multinomial', solver='newton-cg', class_weight="balanced")
#visualizer = ROCAUC(model, classes=['negative', 'positive', 'neutral'])
visualizer = ClassificationReport(model, classes=['negative', 'positive', 'neutral'])

#confusion matrix
log_cm = ConfusionMatrix(model)
log_cm.fit(x_train, y_train)          # Fit the training data to the visualizer
log_cm.score(x_test, y_test)         # Evaluate the model on the test data
log_cm.show()


##classification report
log_cr = ClassificationReport(model, classes=['negative', 'positive', 'neutral'], support=True)
log_cr.fit(x_train, y_train)          # Fit the training data to the visualizer
log_cr.score(x_test, y_test)         # Evaluate the model on the test data
log_cr.show()

##ROC curve
ROC = ROCAUC(model, classes=['negative', 'positive', 'neutral'])
ROC.fit(x_train, y_train)             # Fit the training data to the visualizer
ROC.score(x_test, y_test)             # Evaluate the model on the test data
ROC.show()
```


```
C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:
213: FutureWarning: From version 0.24, get_params will raise an AttributeError
r if a parameter cannot be retrieved as an instance attribute. Previously it
would return None.
FutureWarning)
```



```
C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:
213: FutureWarning: From version 0.24, get_params will raise an AttributeError
r if a parameter cannot be retrieved as an instance attribute. Previously it
would return None.
FutureWarning)
```




C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.
FutureWarning)



C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.

FutureWarning)



Out[97]: <matplotlib.axes._subplots.AxesSubplot at 0x204962902e8>

```
In [149]: y_pred_lr=model.predict(x_test)
          print("Accuracy_logreg:",metrics.accuracy_score(y_test,y_pred_lr))
```

Accuracy_logreg: 0.9187061175777678

2. Multinomial Naive Bayes

```
In [95]: from sklearn.naive_bayes import MultinomialNB
          mnb = MultinomialNB().fit(x_train, y_train)
          y_predmnb = mnb.predict(x_test)
```

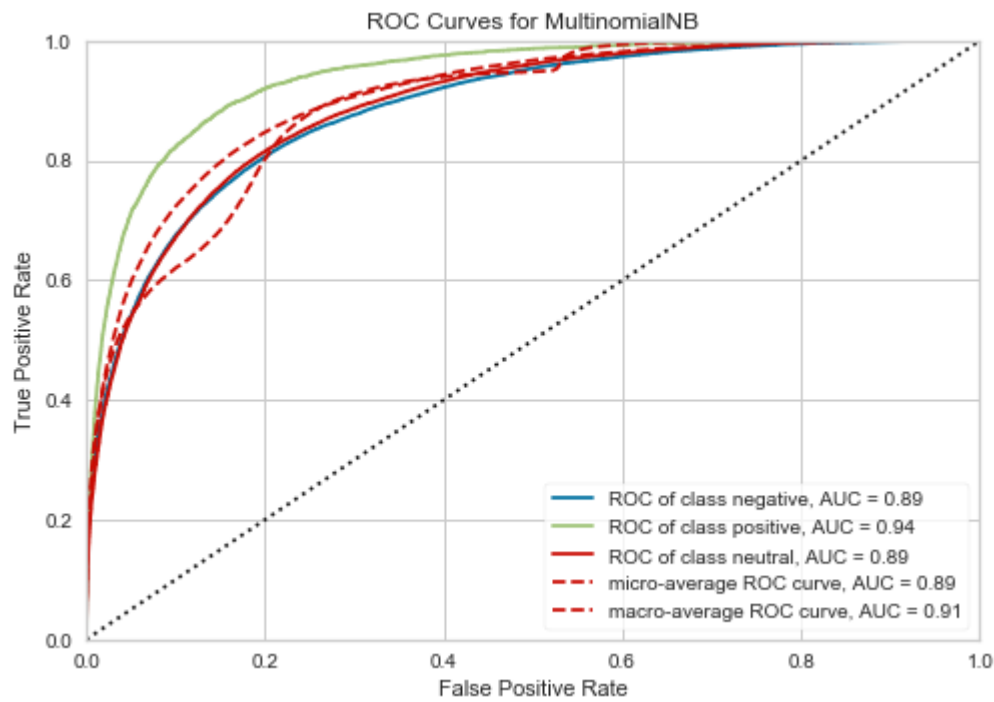
```
In [96]: mnb_roc = ROCAUC(mnb, classes=['negative','positive','neutral'])
#ROC
mnb_roc.fit(x_train, y_train)           # Fit the training data to the visualizer
mnb_roc.score(x_train, y_train)         # Evaluate the model on the test data
mnb_roc.show()

#confusion matrix
mnb_cm = ConfusionMatrix(mnb)
mnb_cm.fit(x_train, y_train)           # Fit the training data to the visualizer
mnb_cm.score(x_test, y_test)          # Evaluate the model on the test data
mnb_cm.show()


##classification report

mnb_cr = ClassificationReport(mnb, classes=['negative','positive','neutral'],
support=True)
mnb_cr.fit(x_train, y_train)           # Fit the training data to the visualizer
mnb_cr.score(x_test, y_test)          # Evaluate the model on the test data
mnb_cr.show()
```

```
C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:
213: FutureWarning: From version 0.24, get_params will raise an AttributeError
r if a parameter cannot be retrieved as an instance attribute. Previously it
would return None.
FutureWarning)
```




```
C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:
213: FutureWarning: From version 0.24, get_params will raise an AttributeError
r if a parameter cannot be retrieved as an instance attribute. Previously it
would return None.
FutureWarning)
```



C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.

FutureWarning)



Out[96]: <matplotlib.axes._subplots.AxesSubplot at 0x2049387b128>

```
In [153]: #y_pred_lr=model.predict(x_test)
print("Accuracy_navibayes:",metrics.accuracy_score(y_test,y_predmnmb))
```

Accuracy_navibayes: 0.6723755792867218

3. Decision Tree Classifier

```
In [99]: import sklearn.metrics as metrics
from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt = dt.fit(x_train,y_train)

#Predict the response for test dataset
y_predtree = dt.predict(x_test)
print("Accuracy:tree",metrics.accuracy_score(y_test, y_predtree))
```

Accuracy:tree 0.9020288597157426

```
In [100]: dt_roc = ROCAUC(dt, classes=['negative','positive','neutral'])
dt_roc.fit(x_train, y_train)           # Fit the training data to the visualizer
dt_roc.score(x_test, y_test)          # Evaluate the model on the test data
dt_roc.show()


#confusion matrix
dt_cm = ConfusionMatrix(dt)
dt_cm.fit(x_train, y_train)           # Fit the training data to the visualizer
dt_cm.score(x_test, y_test)          # Evaluate the model on the test data
dt_cm.show()

##classification report


dt_cr = ClassificationReport(dt, classes=['negative','positive','neutral'], support=True)
dt_cr.fit(x_train, y_train)           # Fit the training data to the visualizer
dt_cr.score(x_test, y_test)          # Evaluate the model on the test data
dt_cr.show()
```



```
C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:
213: FutureWarning: From version 0.24, get_params will raise an AttributeError
r if a parameter cannot be retrieved as an instance attribute. Previously it
would return None.
FutureWarning)
```




```
C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:
213: FutureWarning: From version 0.24, get_params will raise an AttributeError
r if a parameter cannot be retrieved as an instance attribute. Previously it
would return None.
FutureWarning)
```



C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.

FutureWarning)



Out[100]: <matplotlib.axes._subplots.AxesSubplot at 0x204d872de48>

In [154]: `print("Accuracy_Decisiontree:", metrics.accuracy_score(y_test, y_predtree))`

Accuracy_Decisiontree: 0.9020288597157426

4. Random Forest

In [102]: `from sklearn.ensemble import RandomForestClassifier`
`rf = RandomForestClassifier(n_estimators=100)`

```
In [106]: rf_roc = ROCAUC(rf, classes=['negative','positive','neutral'])
rf_roc.fit(x_train, y_train)           # Fit the training data to the visualizer
rf_roc.score(x_test, y_test)          # Evaluate the model on the test data
rf_roc.show()


#confusion matrix
rf_cm = ConfusionMatrix(rf)
rf_cm.fit(x_train, y_train)           # Fit the training data to the visualizer
rf_cm.score(x_test, y_test)          # Evaluate the model on the test data
rf_cm.show()

##classification report

rf_cr = ClassificationReport(rf, classes=['negative','positive','neutral'], support=True)
rf_cr.fit(x_train, y_train)           # Fit the training data to the visualizer
rf_cr.score(x_test, y_test)          # Evaluate the model on the test data
rf_cr.show()
```


C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.

FutureWarning)




C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.

FutureWarning)



C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.

FutureWarning)



Out[106]: <matplotlib.axes._subplots.AxesSubplot at 0x204404a4358>

```
In [155]: y_pred_rf=rf.predict(x_test)
print("Accuracy_randomforest:",metrics.accuracy_score(y_test,y_pred_rf))
```

Accuracy_randomforest: 0.8956121452595359

5. Support Vector Machines

```
In [107]: from sklearn import svm
clf = svm.SVC(kernel='linear',probability = True) # Linear Kernel

#Train the model using the training sets
svm = clf.fit(x_train, y_train)
```

```
In [108]: #Predict the response for test dataset
y_pred_svm = clf.predict(x_test)
```


```
In [117]: print("Accuracy_svm:",metrics.accuracy_score(y_test,y_pred_svm))
```

Accuracy_svm: 0.9501697174475736

```
In [110]: svm_roc = ROCAUC(svm, classes=['negative','positive','neutral'])
svm_roc.fit(x_train, y_train)           # Fit the training data to the visualizer
svm_roc.score(x_test, y_test)           # Evaluate the model on the test data
svm_roc.show()
```

C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:
213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.


FutureWarning)



Out[110]: <matplotlib.axes._subplots.AxesSubplot at 0x204cba53438>

```
In [115]: svm_cr = ClassificationReport(svm, classes=['negative', 'positive', 'neutral'],  
support=True)  
svm_cr.score(x_test, y_test)      # Evaluate the model on the test data  
svm_cr.show()
```


C:\ProgramData\Anaconda3\lib\site-packages\yellowbrick\classifier\base.py:23
2: YellowbrickWarning: could not determine class_counts_ from previously fitted classifier
YellowbrickWarning,



Out[115]: <matplotlib.axes._subplots.AxesSubplot at 0x2043f130748>

```
In [118]: #confusion matrix
svm_cm = ConfusionMatrix(svm)
svm_cm.fit(x_train, y_train) # Fit the training data to the visualizer
svm_cm.score(x_test, y_test) # Evaluate the model on the test data
svm_cm.show()
```

C:\Users\adari\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py:213: FutureWarning: From version 0.24, get_params will raise an AttributeError if a parameter cannot be retrieved as an instance attribute. Previously it would return None.
FutureWarning)



Out[118]: <matplotlib.axes._subplots.AxesSubplot at 0x204cc008a90>

Feature importance

```
In [136]: from yellowbrick.model_selection import FeatureImportances
```

```
In [139]: feature_importances = pd.DataFrame(dt.feature_importances_,
                                             columns=['importance']).sort_value
s('importance', ascending=False)
```



```
In [141]: feature_importances.head(5)
```

```
Out[141]:
```

	importance
16449	0.044340
39901	0.042791
37404	0.030712
84651	0.028394
37083	0.027371

```
In [140]: feature_imp_rf = pd.DataFrame(rf.feature_importances_,  
                                         columns=['importance']).sort_value  
s('importance', ascending=False)
```

```
In [143]: feature_imp_rf.head(6)
```

```
Out[143]:
```

	importance
16449	0.014443
39901	0.011668
37404	0.011634
84651	0.010440
37083	0.008879
76637	0.008822

```
In [145]: wordtf = pd.DataFrame(v.get_feature_names(), columns = ['wrd'])
```

```
In [147]: wordtf.iloc[[16449,39901,37404,84651,37083,76637],]
```

```
Out[147]:
```

	wrd
16449	bad
39901	horrible
37404	great
84651	worst
37083	good
76637	terrible

```
In [ ]:
```